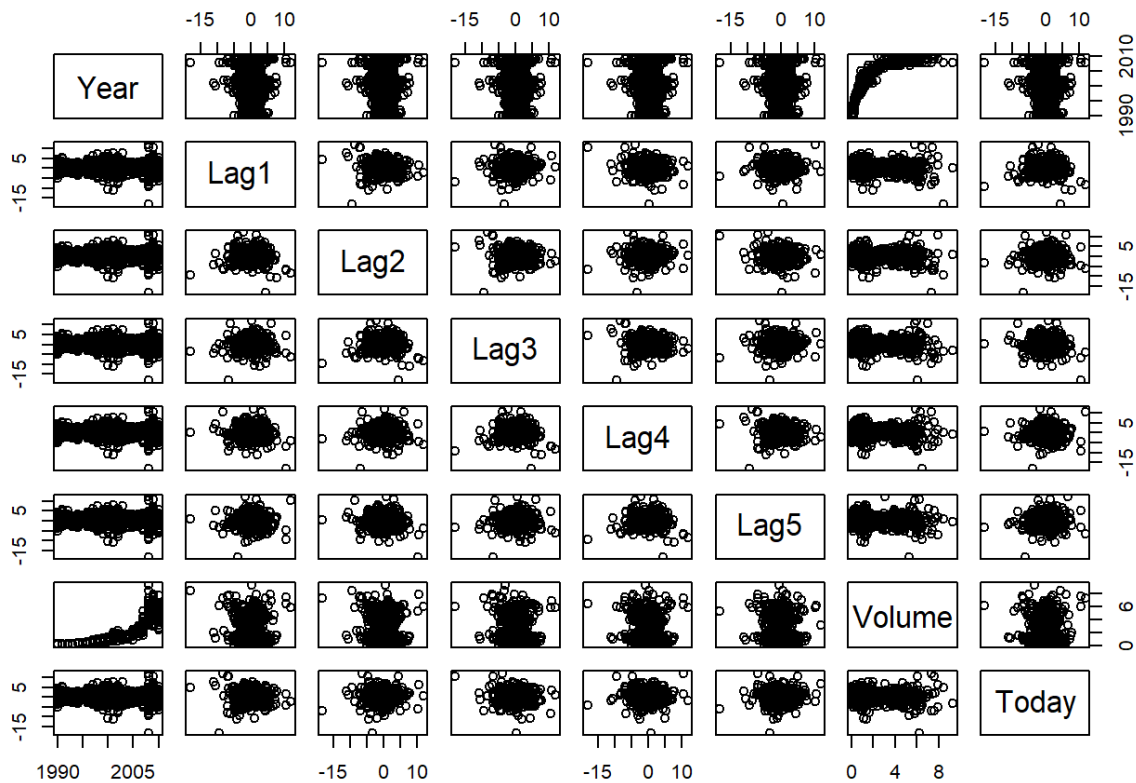


Question1-

a-

```
pairs(Weekly[, -9])
```



```
abs(cor(Weekly[, -9]))
```

##	Year	Lag1	Lag2	Lag3	Lag4	Lag5
## Year	1.00000000	0.032289274	0.03339001	0.03000649	0.031127923	0.03051910
## Lag1	0.03228927	1.00000000	0.07485305	0.05863568	0.071273876	0.00818309
## Lag2	0.03339001	0.074853051	1.00000000	0.07572091	0.058381535	0.07249948
## Lag3	0.03000649	0.058635682	0.07572091	1.00000000	0.075395865	0.06065717
## Lag4	0.03112792	0.071273876	0.05838153	0.07539587	1.00000000	0.07567502
## Lag5	0.03051910	0.00818309	0.07249948	0.06065717	0.07567502	1.00000000

Reihaneh Moghisi

Assignment 2

RMI8300


```
## Lag5    0.03051910  0.008183096  0.07249948  0.06065717  0.075675027  1.00000000
0

## Volume  0.84194162  0.064951313  0.08551314  0.06928771  0.061074617  0.05851741
4

## Today   0.03245989  0.075031842  0.05916672  0.07124364  0.007825873  0.01101269
8

##          Volume          Today
## Year    0.84194162  0.032459894
## Lag1     0.06495131  0.075031842
## Lag2     0.08551314  0.059166717
## Lag3     0.06928771  0.071243639
## Lag4     0.06107462  0.007825873
## Lag5     0.05851741  0.011012698
## Volume   1.00000000  0.033077783
## Today    0.03307778  1.000000000
```

There are no obvious strong relationships between the `Lag` variables. Although there is some trends of volume over time.



B — only Lag 2 is statistically significant.

Reihaneh Moghisi

Assignment 2

RMI8300

```
glm_dir <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,
               data = Weekly,
               family = "binomial")
```

```
summary(glm_dir)
```

```
##
```

```
## Call:
```

```
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
```

```
##      Volume, family = "binomial", data = Weekly)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -1.6949  -1.2565   0.9913   1.0849   1.4579
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.26686    0.08593   3.106   0.0019 **
## Lag1         -0.04127    0.02641  -1.563   0.1181
## Lag2          0.05844    0.02686   2.175   0.0296 *
## Lag3         -0.01606    0.02666  -0.602   0.5469
## Lag4         -0.02779    0.02646  -1.050   0.2937
## Lag5         -0.01447    0.02638  -0.549   0.5833
## Volume       -0.02274    0.03690  -0.616   0.5377
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
```

```
##      Null deviance: 1496.2  on 1088  degrees of freedom
```

```
## Residual deviance: 1486.4  on 1082  degrees of freedom
```

```
## AIC: 1500.4
```

```
##
```

C –

The confusion matrix is shown below. The confusion matrix in caret package is used for this. It shows a poor accuracy of 56% . this prediction is based on only training set. It also shows very low specificity which means did not perform well in predicting negative class correctly.

```
predicted <- factor(ifelse(predict(glm_dir, type = "response") < 0.5, "Down",  
"Up"))
```

```
confusionMatrix(predicted, Weekly$Direction, positive = "Up")
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction Down  Up
```

```
##           Down   54  48
```

```
##           Up    430 557
```

```
##
```

```
##           Accuracy : 0.5611
```

```
##           95% CI : (0.531, 0.5908)
```

```
## No Information Rate : 0.5556
```

```
## P-Value [Acc > NIR] : 0.369
```

```
##
```

```
##           Kappa : 0.035
```

```
##
```

```
## McNemar's Test P-Value : <2e-16
```

```
##
```

```
##           Sensitivity : 0.9207
```

```
##           Specificity : 0.1116
```

```
## Pos Pred Value : 0.5643
```

```
## Neg Pred Value : 0.5294
```

```
##           Prevalence : 0.5556
```

```
## Detection Rate : 0.5115
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
##      Detection Prevalence : 0.9063
##      Balanced Accuracy   : 0.5161
##
##      'Positive' Class    : Up
##
```

D –

The accuracy of model is **0.625**.

No Information Rate : 0.5865 tells us that the largest class (Up) is 58.65% of the test observations, this is the baseline. As the Accuracy : 0.625 > 0.5865 shows a positive outcome of prediction however because the test dataset is relatively small so this might not be a meaningful improvement.

The p-value for P-Value [Acc > NIR] : 0.2439 is a one-sided test to see whether the accuracy is better than the “no information rate. And as this p-value is greater than > 0.05 we conclude that there is no significant evidence that our classifier is better than the baseline strategy.

```
train <- Weekly[Weekly$Year <= 2008, ]
test  <- Weekly[Weekly$Year > 2008, ]

glm_dir <- glm(Direction ~ Lag2,
               data = train,
               family = "binomial")

predicted <- factor(ifelse(predict(glm_dir, newdata = test, type = "response")
                             < 0.5, "Down", "Up"))

confusionMatrix(predicted, test$Direction, positive = "Up")

## Confusion Matrix and Statistics
##
##      Reference
## Prediction Down Up
##      Down      9   5
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
##           Up           34  56
##
##           Accuracy : 0.625
##           95% CI : (0.5247, 0.718)
##           No Information Rate : 0.5865
##           P-Value [Acc > NIR] : 0.2439
##
##           Kappa : 0.1414
##
##  Mcnemar's Test P-Value : 7.34e-06
##
##           Sensitivity : 0.9180
##           Specificity : 0.2093
##           Pos Pred Value : 0.6222
##           Neg Pred Value : 0.6429
##           Prevalence : 0.5865
##           Detection Rate : 0.5385
##           Detection Prevalence : 0.8654
##           Balanced Accuracy : 0.5637
##
##           'Positive' Class : Up
##
```

E – Repeat (d) using LDA

Confusion Matrix and Statistics

```
##
##           Reference
## Prediction Down Up
##           Down      9   5
##           Up       34  56
##
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
##                      Accuracy : 0.625
##                      95% CI : (0.5247, 0.718)
##      No Information Rate : 0.5865
##      P-Value [Acc > NIR] : 0.2439
##
##                      Kappa : 0.1414
##
##      McNemar's Test P-Value : 7.34e-06
##
##                      Sensitivity : 0.9180
##                      Specificity : 0.2093
##                      Pos Pred Value : 0.6222
##                      Neg Pred Value : 0.6429
##                      Prevalence : 0.5865
##                      Detection Rate : 0.5385
##      Detection Prevalence : 0.8654
##      Balanced Accuracy : 0.5637
##
##                      'Positive' Class : Up
##
```

The Accuracy of model is **0.625**. Same as above, the P-Value [Acc > NIR] : 0.2439 > 0.05, so although the accuracy of the classifier is 0.625 > 0.5865, with the same logic that the test sample size is not large enough and therefore the increase in accuracy is not statistically significant enough over the baseline.

F- Repeat using QDA.

```
qda_dir <- qda(Direction ~ Lag2, data = train)

predicted_qda <- predict(qda_dir, newdata = test)
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
confusionMatrix(data = predicted_gda$class,
                 reference = test$Direction,
                 positive = "Up")

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Down Up
##      Down      0   0
##      Up       43  61
##
##           Accuracy : 0.5865
##           95% CI : (0.4858, 0.6823)
##      No Information Rate : 0.5865
##      P-Value [Acc > NIR] : 0.5419
##
##           Kappa : 0
##
##  Mcnemar's Test P-Value : 1.504e-10
##
##           Sensitivity : 1.0000
##           Specificity : 0.0000
##      Pos Pred Value : 0.5865
##      Neg Pred Value :      NaN
##           Prevalence : 0.5865
##      Detection Rate : 0.5865
##      Detection Prevalence : 1.0000
##      Balanced Accuracy : 0.5000
##
##      'Positive' Class : Up
##
```

Here we get an Accuracy of **0.5865**. with no-information rate of 58% and p-value 0.54. the accuracy and no-information rate are equal and specificity 0 and sensitivity 1 illustrate that this model follows naïve classifier.

G- Repeat using KNN

the confusion matrix:

```
confusionMatrix(data = predicted_knn,
                 reference = test$Direction,
                 positive = "Up")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Down Up
##      Down   21 30
##      Up    22 31
##
##              Accuracy : 0.5
##              95% CI : (0.4003, 0.5997)
##      No Information Rate : 0.5865
##      P-Value [Acc > NIR] : 0.9700
##
##              Kappa : -0.0033
##
##      McNemar's Test P-Value : 0.3317
##
##              Sensitivity : 0.5082
##              Specificity : 0.4884
##      Pos Pred Value : 0.5849
##      Neg Pred Value : 0.4118
##              Prevalence : 0.5865
##      Detection Rate : 0.2981
##      Detection Prevalence : 0.5096
##      Balanced Accuracy : 0.4983
##
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
##           'Positive' Class : Up
##
```

Here we get an accuracy of **0.5**, which is again worse than the baseline.

H- Which of these methods appears to provide the best results on this data?

Considering the accuracy as the measurement of performance: LDA & Logistic Regression get the same test accuracy of **0.625**, so these two are best classifiers in this case.

I – combined predictors, interactions, transformations

```
ctrl <- trainControl(method = "repeatedcv",
                     number = 5,
                     repeats = 5)

set.seed(111)

knn_train <- train(y = train$Direction,
                  x = train[, -8],
                  method = "knn",
                  metric = "Accuracy",
                  preProcess = c("center", "scale"),
                  tuneGrid = expand.grid(k = seq(1, 50, 2)),
                  trControl = ctrl)

caret::varImp(knn_train)

## ROC curve variable importance
##
##      Importance
## Lag1      100.000
## Lag2       77.256
## Lag5       64.309
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
## Year      45.659
## Volume    43.735
## Week      42.513
## Lag4       4.578
## Lag3       0.000
```

```
knn_train
## k-Nearest Neighbors
##
## 985 samples
##    8 predictor
##    2 classes: 'Down', 'Up'
##
## Pre-processing: centered (8), scaled (8)
## Resampling: Cross-Validated (5 fold, repeated 5 times)
## Summary of sample sizes: 788, 788, 788, 788, 788, 787, ...
## Resampling results across tuning parameters:
##
##    k    Accuracy    Kappa
##    1  0.4947996  -0.020756148
##    3  0.5232477   0.031373990
##    5  0.5230292   0.028769372
##    7  0.5372435   0.053353378
##    9  0.5372415   0.049612758
##   11  0.5343834   0.040819116
##   13  0.5346081   0.037598994
##   15  0.5368437   0.040363712
##   17  0.5317675   0.026161027
##   19  0.5301555   0.021608137
##   21  0.5370622   0.033488872
##   23  0.5330064   0.024861594
##   25  0.5401182   0.036317004
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
##      27  0.5376693  0.029998090
##      29  0.5321890  0.015868622
##      31  0.5338310  0.018687256
##      33  0.5360563  0.021431639
##      35  0.5322138  0.012669022
##      37  0.5283538  0.002535516
##      39  0.5330239  0.011432961
##      41  0.5289589  0.002400751
##      43  0.5295618  0.003081904
##      45  0.5299700  0.003587628
##      47  0.5303854  0.003745293
##      49  0.5320056  0.005552184
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 25.
```

Evaluating the performance of this new model on test:

```
knn_pred <- predict(knn_train, newdata = test)

confusionMatrix(data = knn_pred,
                 reference = test$Direction,
                 positive = "Up")

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Down Up
##      Down   19 20
##      Up    24 41
##
##              Accuracy : 0.5769
##              95% CI : (0.4761, 0.6732)
##      No Information Rate : 0.5865
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
##      P-Value [Acc > NIR] : 0.6193
##
##      Kappa : 0.1156
##
##      McNemar's Test P-Value : 0.6511
##
##      Sensitivity : 0.6721
##      Specificity : 0.4419
##      Pos Pred Value : 0.6308
##      Neg Pred Value : 0.4872
##      Prevalence : 0.5865
##      Detection Rate : 0.3942
##      Detection Prevalence : 0.6250
##      Balanced Accuracy : 0.5570
##
##      'Positive' Class : Up
##
```

Although the accuracy increased from k=1 but we are still below the baseline.

Another combination is using the interaction term of Lag1 and Lag2. The results below shows still very poor accuracy of 55% over baseline.

```
glm(formula = Direction ~ Lag1 * Lag2 + Lag3 + Lag4 + Lag5 +
     Volume, family = "binomial", data = Weekly)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6176	-1.2590	0.9893	1.0852	1.5367

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.268646	0.086017	3.123	0.00179 **
Lag1	-0.036253	0.028227	-1.284	0.19903
Lag2	0.057800	0.027025	2.139	0.03246 *
Lag3	-0.017376	0.026826	-0.648	0.51716
Lag4	-0.029814	0.026775	-1.113	0.26550
Lag5	-0.014385	0.026380	-0.545	0.58553
Volume	-0.023317	0.036951	-0.631	0.52802
Lag1:Lag2	0.003475	0.006891	0.504	0.61410

(Dispersion parameter for binomial family taken to be 1)

Reihaneh Moghisi

Assignment 2

RMI8300

Null deviance: 1496.2 on 1088 degrees of freedom
Residual deviance: 1486.1 on 1081 degrees of freedom
AIC: 1502.1

Number of Fisher Scoring iterations: 4

Confusion Matrix and Statistics

	Reference	
Prediction	Down	Up
Down	51	48
Up	433	557

Accuracy : 0.5583
95% CI : (0.5282, 0.5881)
No Information Rate : 0.5556
P-Value [Acc > NIR] : 0.4399

Kappa : 0.0283

McNemar's Test P-Value : <2e-16

Sensitivity : 0.9207
Specificity : 0.1054
Pos Pred Value : 0.5626
Neg Pred Value : 0.5152
Prevalence : 0.5556
Detection Rate : 0.5115
Detection Prevalence : 0.9091
Balanced Accuracy : 0.5130

'Positive' Class : Up

Question 2-

A –

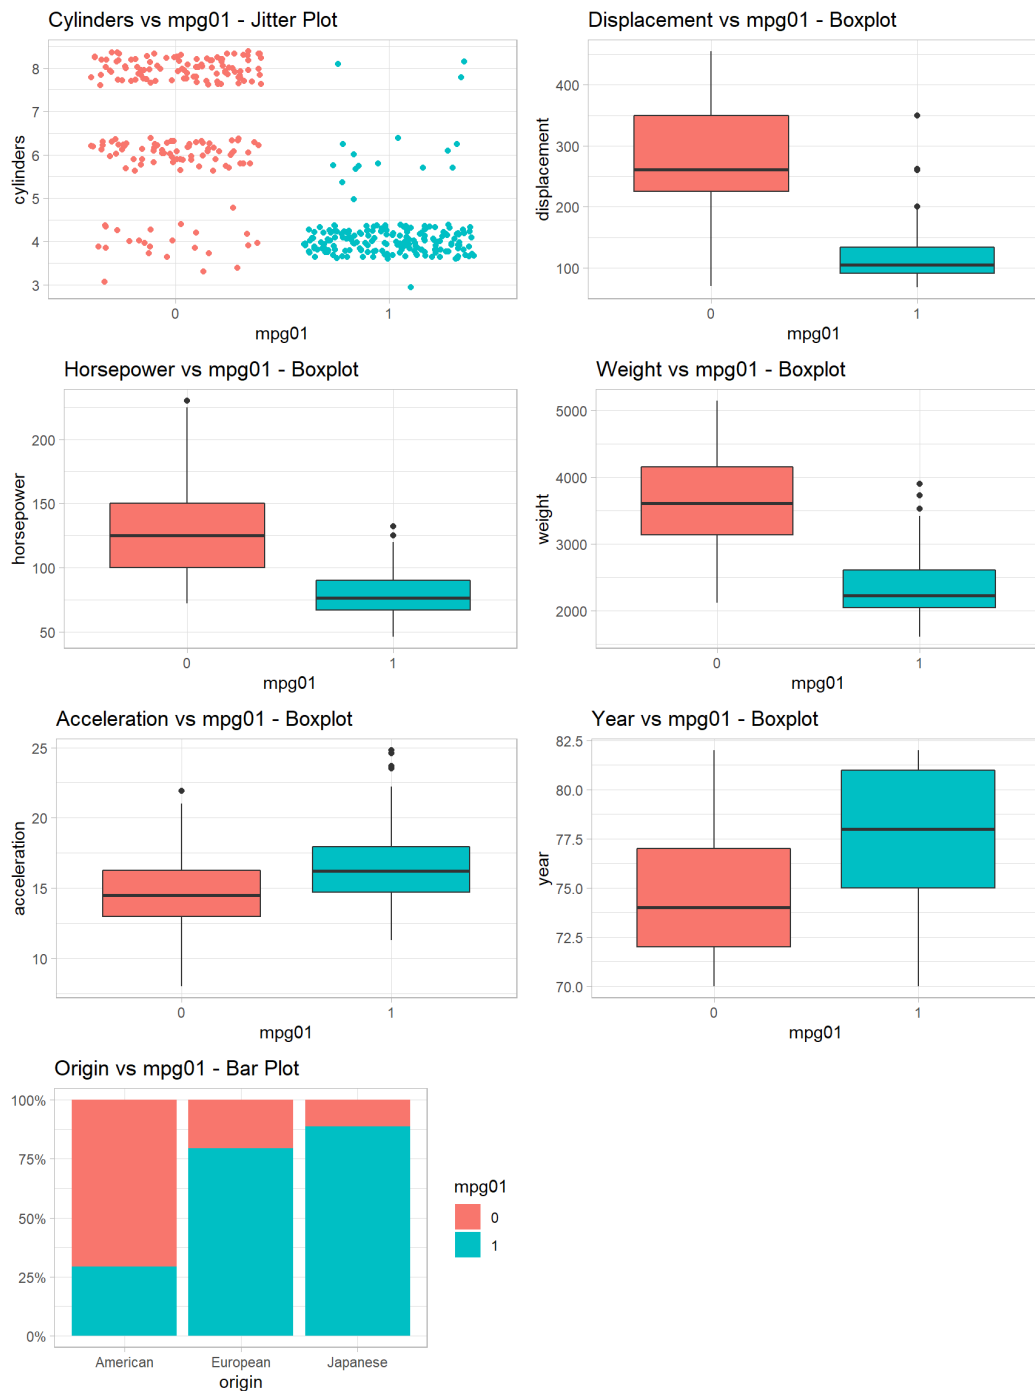
```
Auto$mpg01 <- factor(as.numeric(Auto$mpg > median(Auto$mpg)))  
table(Auto$mpg01)  
##  
##    0    1  
## 196 196
```

Reihaneh Moghisi

Assignment 2

RMI8300

B — Although the exact associations can't be extracted from graphical plots but we can poorly conclude that `cylinders`, `displacement` & `weight` are the three strongest predictors of `mpg01`. Horsepower, year and acceleration look weaker predictors of `mpg01`.



C-

```
set.seed(444)

index <- createDataPartition(y = Auto$mpg01, p = 0.5, list = F)

train <- Auto[index, ]
test <- Auto[-index, ]

nrow(train) / nrow(Auto)

## [1] 0.5
```

D-

We use the top 4 strongest predictors I identified visually
(cylinders, displacement, weight & horsepower).

It shows that the model is doing great in predicting the dependent variable much better than baseline.

The test error is 0.12244

```
ctrl <- trainControl(method = "repeatedcv", number = 10, repeats = 3)

set.seed(1)

lda_mpg <- train(mpg01 ~ cylinders + displacement + weight + horsepower,
                  data = train,
                  method = "lda",
                  trControl = ctrl)
```

train (cross-validation) error:

```
1 - lda_mpg$results$Accuracy # cv error

## [1] 0.09125731
```

test error:

Reihaneh Moghisi

Assignment 2

RMI8300

```
predicted_lda <- predict(lda_mpg, newdata = test, type = "raw") # as opposed
to type = "prob"

mean(predicted_lda != test$mpg01)

## [1] 0.122449
```

E-

In this case, QDA appears to perform better than LDA with respect to test error, and slightly better in terms of CV error.

```
set.seed(2)

qda_mpg <- train(mpg01 ~ cylinders + displacement + weight + horsepower,
                 data = train,
                 method = "qda",
                 trControl = ctrl)
```

train (cross-validation) error:

```
1 - qda_mpg$results$Accuracy

## [1] 0.08991228
```

test error:

```
predicted_qda <- predict(qda_mpg, newdata = test, type = "raw")

mean(predicted_qda != test$mpg01)

## [1] 0.1173469
```

F-

We can see here that Logistic Regression performs better than LDA & QDA with respect to CV error & test error.

```
set.seed(3)
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
log_mpg <- train(mpg01 ~ cylinders + displacement + weight + horsepower,
                 data = train,
                 method = "glm",
                 family = "binomial",
                 trControl = ctrl)
```

train (cross-validation) error:

```
1 - log_mpg$results$Accuracy
## [1] 0.1084405
```

test error:

```
predicted_log <- predict(log_mpg, newdata = test, type = "raw")

mean(predicted_log != test$mpg01)
## [1] 0.1173469
```

G- The optimal value chosen was $K = 7$, although this model performed close to Logistic Regression in CV error, and slightly worse in test error.

It is observed that all classifiers with larger values for K (from 10 to 73) get the same cross-validation accuracy scores. This can be due to the fact that those KNN models make the same prediction.

```
set.seed(4)

knn_mpg <- train(mpg01 ~ cylinders + displacement + weight + horsepower,
                 data = train,
                 method = "knn",
                 preProcess = c("center", "scale"),
                 trControl = ctrl,
                 tuneGrid = expand.grid(k = seq(1, 85, 3)))

knn_mpg
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
## k-Nearest Neighbors
```

```
##
```

```
## 196 samples
```

```
## 4 predictor
```

```
## 2 classes: '0', '1'
```

```
##
```

```
## Pre-processing: centered (4), scaled (4)
```

```
## Resampling: Cross-Validated (10 fold, repeated 3 times)
```

```
## Summary of sample sizes: 176, 177, 176, 176, 176, 176, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

##	k	Accuracy	Kappa
##	1	0.8779532	0.7554422
##	4	0.8999805	0.7995964
##	7	0.9033138	0.8062631
##	10	0.9033138	0.8062631
##	13	0.9033138	0.8062631
##	16	0.9033138	0.8062631
##	19	0.9033138	0.8062631
##	22	0.9033138	0.8062631
##	25	0.9033138	0.8062631
##	28	0.9033138	0.8062631
##	31	0.9033138	0.8062631
##	34	0.9033138	0.8062631
##	37	0.9033138	0.8062631
##	40	0.9033138	0.8062631
##	43	0.9033138	0.8062631
##	46	0.9033138	0.8062631
##	49	0.9033138	0.8062631
##	52	0.9033138	0.8062631
##	55	0.9033138	0.8062631
##	58	0.9033138	0.8062631
##	61	0.9033138	0.8062631

Reihaneh Moghisi

Assignment 2

RMI8300

```
##    64    0.9049805    0.8095964
##    67    0.9049805    0.8095964
##    70    0.9066472    0.8129297
##    73    0.9100682    0.8197621
##    76    0.9083138    0.8162631
##    79    0.9084016    0.8164288
##    82    0.9067349    0.8130955
##    85    0.9034016    0.8065072

##

## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 73.
```

train (cross-validation) error:

```
## [1] 0.08993177
```

test error:

```
## [1] 0.1122449
```

Question 3:

A –

```
Power <- function() {
  2^3
}

Power()

## [1] 8
```

B –

```
Power2 <- function(x, a) {
  x^a
}
```

Reihaneh Moghisi
Assignment 2
RMI8300

```
Power2(3, 8)
## [1] 6561
```

C –

```
Power2(10, 3)
## [1] 1000
```

8 17:

```
Power2(8, 17)
## [1] 2.2518e+15
```

131 3:

```
Power2(131, 3)
## [1] 2248091
```


D-

```
Power3 <- function(x, a) {
  result <- x^a
  result
}

Power3(3, 8)
## [1] 6561
```

E –

```
data.frame(x = 1:10, y = Power3(1:10, 2)) %>%
  ggplot(aes(x = x, y = log(y))) +
  geom_point() +
  geom_line() +
  scale_x_continuous(breaks = 1:10, minor_breaks = NULL)
```



F-

```
PlotPower <- function(x, a, col = "black") {  
  ggplot(mapping = aes(x = x, y = x^a)) +  
    geom_point(col = col) +  
    geom_line(col = col) +  
    scale_x_continuous(labels = scales::comma_format()) +  
    scale_y_continuous(labels = scales::comma_format()) +  
    theme_light() +  
    labs(title = paste0("Plot of f(x) = x^", as.character(a), " (for x between  
n ", min(x), " and ", max(x), ")"),  
         subtitle = "Created using the 'PlotPower()' function")  
}
```

I test the `PlotPower()` function for two examples below. I also added a `col` argument to change the colour, and made the title auto-generate to display the function and range.

Reihaneh Moghisi


Assignment 2

RMI8300

`PlotPower(1:20, 3)`

Plot of $f(x) = x^3$ (for x between 1 and 20)


Created using the 'PlotPower()' function



```
PlotPower(x = seq(1, 500, 10),  
          a = 2,  
          col = "deepskyblue3")
```

Plot of $f(x) = x^2$ (for x between 1 and 491)

Created using the 'PlotPower()' function



Question 4:

Reihaneh Moghisi

Assignment 2

RM18300

First, I'll overwrite the crime rate variable (`crim`) with a binary factor variable, which takes the value 1 when `crim` is above the median, else 0.

Then I start by fitting all 4 model types, using all predictors each time.

Instead of `train/test` split, I used using cross-validation to evaluate performance of accuracy (10-fold, repeated 5 times).

From the results we can see that KNN performed the best. Logistic regression was the second highest accuracy followed by QDA and LDA .

Changing the crime variable:

```
Boston$crim <- factor(ifelse(Boston$crim > median(Boston$crim), 1, 0))
```

Cross validation:

```
ctrl <- trainControl(method = "repeatedcv",  
                     number = 10,  
                     repeats = 5)
```

Logistic Regression:

```
set.seed(1)  
log_crim <- train(crim ~ .,  
                 data = Boston,  
                 method = "glm",  
                 family = "binomial",  
                 trControl = ctrl)
```

```
log_crim  
## Generalized Linear Model  
##  
## 506 samples  
## 13 predictor  
## 2 classes: '0', '1'
```


Reihaneh Moghisi

Assignment 2

RMI8300

```
##  
## No pre-processing  
## Resampling: Cross-Validated (10 fold, repeated 5 times)  
## Summary of sample sizes: 456, 455, 456, 454, 456, 456, ...  
## Resampling results:  
##  
##   Accuracy   Kappa  
##   0.9028549  0.8057423
```

LDA:

```
set.seed(2)  
  
lda_crim <- train(crim ~ .,  
                  data = Boston,  
                  method = "lda",  
                  trControl = ctrl)  
  
lda_crim  
## Linear Discriminant Analysis  
##  
## 506 samples  
## 13 predictor  
## 2 classes: '0', '1'  
##  
## No pre-processing  
## Resampling: Cross-Validated (10 fold, repeated 5 times)  
## Summary of sample sizes: 456, 455, 456, 456, 455, 454, ...  
## Resampling results:  
##  
##   Accuracy   Kappa  
##   0.8498998  0.6996644
```

QDA:

```
set.seed(3)

qda_crim <- train(crim ~ .,
                  data = Boston,
                  method = "qda",
                  trControl = ctrl)

qda_crim

## Quadratic Discriminant Analysis
##
## 506 samples
## 13 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 456, 456, 455, 455, 455, 456, ...
## Resampling results:
##
## Accuracy   Kappa
## 0.8894561  0.7788975
```

KNN:

```
set.seed(4)

knn_crim <- train(crim ~ .,
                  data = Boston,
                  method = "knn",
                  preProcess = c("center", "scale"),
```

Reihaneh Moghisi

Assignment 2

RMI8300

```
trControl = ctrl,
tuneGrid = expand.grid(k = seq(1, 20, 2))

knn_crim
## k-Nearest Neighbors
##
## 506 samples
## 13 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 455, 456, 456, 455, 455, 456, ...
## Resampling results across tuning parameters:
##
##    k    Accuracy    Kappa
##    1  0.9145970  0.8292052
##    3  0.9181104  0.8362439
##    5  0.9153345  0.8306748
##    7  0.9047128  0.8094000
##    9  0.8779164  0.7557756
##   11  0.8664404  0.7327928
##   13  0.8664323  0.7327720
##   15  0.8703940  0.7407136
##   17  0.8684169  0.7367859
##   19  0.8648404  0.7296460
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 3.
```